

MMI713 Term Project - Phase 3 Report

Abdullah Emir Göğüsdere
Electrical and Electronics Engineering
Middle East Technical University
Ankara, Turkey
e244307@metu.edu.tr

Abstract—Solving large sparse linear systems efficiently is a fundamental challenge in scientific computing, particularly in fields like circuit simulation, structural analysis, and machine learning. The Conjugate Gradient (CG) algorithm is widely used for such problems, but its performance can be limited by sequential bottlenecks and memory-bound operations on traditional CPUs. This project presents a GPU-accelerated implementation of the CG method using CUDA, featuring a row binning strategy that assigns sparse and dense rows to separate CUDA kernels (spmv_bin0_kernel and spmv_bin1_kernel) for optimized execution. By using concurrent kernel launches via CUDA streams and carefully balancing workloads, the implementation achieves up to 39× speedup over a sequential CPU version on large benchmark matrices. The implementation code, data files (in .txt format), and documentation are available at: [GitHub Repository](#).

Index Terms—Conjugate Gradient, Sparse Linear Systems, GPU Acceleration, CUDA, Sparse Matrix-Vector Multiplication, Stream Concurrency, Row Binning, Parallel Computing

I. MOTIVATION & SIGNIFICANCE

Solving large sparse linear systems efficiently is a fundamental challenge in many scientific and engineering applications. As system sizes and problem complexities increase in scientific and engineering computations, the time spent in communication and data exchange across processors can dominate the total computation time. This makes it inefficient to solve large-scale linear systems. For these systems, the Conjugate Gradient (CG) algorithm is one of the most widely used iterative solvers due to its low memory requirements and strong convergence properties [1].

However, the performance of CG is heavily dependent on the efficiency of underlying linear algebra operations, particularly sparse matrix-vector multiplication (SpMV) and inner product calculations. With the increasing size and complexity of modern simulations, traditional CPU-based implementations often become a performance bottleneck. Usage of GPU acceleration provides an opportunity to significantly reduce solution time through parallelism [1].

This work focuses on designing and implementing parallelized CG algorithm optimized for sparse matrices in compressed sparse row (CSR) format on NVIDIA

GPUs. It uses strategies such as row binning, warp-level parallelism, and concurrent kernel execution to improve scalability and throughput. The ability to solve large-scale systems faster and more efficiently is crucial in enabling real-time simulation and faster iterative design in critical domains.

This report includes material from my earlier CENG577 coursework. Some parts are reused directly and all reused content is clearly cited to ensure academic honesty.

II. PROBLEM STATEMENT

A. Conjugate Gradient (CG) Algorithm

The CG algorithm is commonly used to find the minimum value of a multi-dimensional function. In the context of solving matrix equations, it solves the equation $Ax = b$ to find x , where A is a symmetric and positive-definite matrix, and b is the right-hand side vector. Here, the goal is to minimize the residual $r = b - Ax$ [1].

During each iteration, the estimated x solution is updated by moving a specific step size α along the search direction vector p . The step size α is computed based on the current values of A and p at that iteration. The search direction p is adjusted based on changes in the norm of the residual, which also serves as a criterion for convergence.

The pseudo-code for sequential CG is provided below [1]:

```
p = b
r = b - Ax
r_old = r
x = x_0
for i = 0 to maxIter do
    alpha = (r · r) / (p · Ap)
    x = x + alpha * p
    r = r - alpha * Ap
    if sqrt(r · r) < tolerance then
        break
    beta = (r · r) / (r_old · r_old)
    p = r + beta * p
    r_old = r
Solution is x
```

III. PRIOR WORK & LIMITATIONS

Several strategies have been developed to address the challenge of efficient parallel implementation of the CG method. These approaches can be roughly categorized into pipelined CG methods and s-step (communication-avoiding) CG methods.

A. Pipelined Conjugate Gradient methods

These methods reduce the effect of global communication in CG by overlapping all-reduce operations, used for dot products, with independent computations like sparse matrix-vector multiplication (SPMV) and preconditioner application. This is accomplished using non-blocking communication primitives that allow the program to start the communication early in the iteration and continue with other tasks that don't depend on its result. As a result, communication latency is hidden behind ongoing computation which is minimizing the idle time of the processor and improving overall throughput [2].

Research has extended the basic pipelined CG method to include "deeper pipelines" (p(l)-CG) [3], where the global all-to-all reduction phase in each CG iteration is overlapped with not just one, but l subsequent SPMV operations. This deeper pipelining is particularly advantageous when the global communication phase becomes the dominant factor limiting scalability, allowing for further performance improvements on systems with many processors. However, pipelining can impact numerical stability due to the reordering of computations and accumulation of rounding errors, potentially increasing the number of iterations required for convergence.

B. s-step Conjugate Gradient methods

s-step CG methods adopt a different strategy by reformulating the CG algorithm to perform a block of s iterations with significantly fewer global communication steps compared to the standard CG algorithm. This reduction in communication is typically achieved by generating a basis for an s -dimensional Krylov subspace within each block of iterations using mostly local operations. By performing multiple SPMVs and other local computations within each s -step block before any global communication occurs, the communication cost is reduced over s iterations, effectively lowering the communication overhead per iteration [4].

A significant challenge associated with this method is maintaining numerical stability. The repeated SPMV operations within each block can lead to the generation of ill-conditioned Krylov bases resulting in potentially causing a loss of accuracy or slower convergence in finite precision arithmetic.

IV. THEORY / ALGORITHM

I propose a parallelized Conjugate Gradient (CG) solver which uses CUDA-based GPU acceleration to efficiently solve large sparse linear systems. This approach aims to improve performance by maximizing memory throughput and exploiting warp-level parallelism.

The block diagram in Fig. 1 illustrates the overall workflow of the proposed solver.




Fig. 1: The block diagram of the overall workflow of the proposed solver.

The process begins with reading the sparse matrix A (in CSR format) and right-hand-side vector b on the CPU. During preprocessing, the rows of the matrix are analyzed and divided into two bins based on row density (Bin0 for sparse rows, Bin1 for denser rows).

All necessary data, including matrix arrays, vectors, and bin indices, are then transferred to the GPU. The parallelized CG algorithm executes iteratively on the GPU where sparse matrix-vector multiplications are computed using two kernels in separate CUDA streams to maximize parallelism.

Each iteration includes vector updates, dot products, and convergence checks. Once the solution vector x is computed, it is transferred back to the CPU for verification by comparing Ax with b . The result is evaluated for accuracy and correctness.

A. Parallelized CG Algorithm

The CG algorithm is inherently well-suited for data-based parallelism, where data and associated operations are distributed among processing units. In contrast, task-based parallelism, where each processor handles a distinct task, is less applicable here. This is evident from the pseudo-code of sequential CG, where most steps rely on the results of previous steps, leaving limited scope for task-based parallelism [1].

The pseudo-code for the parallelized CG algorithm is

provided below:

```

 $p^{device} = b^{device}$ 
 $r^{device} = b^{device}$ 
 $dot_{r}^{device} = r^{device} \cdot r^{device}$ 
 $dot_{r}^{host} = dot_{r}^{device}$ 
 $x^{device} = 0$ 
for  $i = 0$  to  $maxIter$  do
     $(Ap)^{device} = SpMV(A^{device}, p^{device})$ 
     $dot_{pAp}^{device} = p^{device} \cdot (Ap)^{device}$ 
     $dot_{pAp}^{host} = dot_{pAp}^{device}$ 
     $\alpha^{host} = \frac{dot_{r}^{host}}{dot_{pAp}^{host}}$ 
     $x^{device} = x^{device} + \alpha^{host} p^{device}$ 
     $r^{device} = r^{device} - \alpha^{host} (Ap)^{device}$ 
     $dot_r^{device} = r^{device} \cdot r^{device}$ 
     $dot_r^{host} = dot_r^{device}$ 
    if  $\sqrt{dot_r^{host}} < tolerance$  then
        break
     $\beta^{host} = \frac{dot_r^{host}}{dot_{pAp}^{host}}$ 
     $dot_{r}^{host} = dot_r^{host}$ 
     $p^{device} = r^{device} + \beta^{host} p^{device}$ 
     $x^{host} = x^{device}$ 
Solution is  $x^{host}$ 

```

Where:

- the matrix and vector portions residing in GPU memory are denoted by device, in CPU memory are denoted by host.
- \cdot in $r^{device} \cdot r^{device}$ represents the dot product,
- dot_r represents the dot product of r ,
- SpMV represents the sparse matrix-vector multiplication.

B. Matrix storage format

Sparse matrices, where many of their elements are zero, commonly appear in scientific computing. We can use the sparse nature of these matrices to improve performance by storing and working with only the non-zero elements. A simple method for this is called Compressed Sparse Row (CSR) format [1].

CSR is crucial for efficiently storing and processing sparse matrices. It reduces memory usage by storing only non-zero elements and their positions, making it ideal for large, sparse matrices. CSR is especially efficient for operations like sparse matrix-vector multiplication and is widely used in scientific computing due to its compatibility with parallel processing and numerical libraries [1].

CSR represents a matrix using three arrays:

- **Values:** Non-zero elements row-wise.
- **Column Indices:** Column positions of the non-zero elements.
- **Row Pointers:** Start and end of each row in Values array.

Due to these advantages, all matrices in the implementation are stored in CSR format.

C. Parallelization

In the implementation, parallelization of CG operations is achieved primarily through CUDA-based kernel execution.

1) *SpMV - sparse matrix-vector multiplication:* The most computationally intensive operation, sparse matrix-vector multiplication (SpMV), is optimized by separating matrix rows based on total number of nonzero in a row to enable load balancing. The management of this separation is handled by a method called row binning.

Row binning analyzes the number of nonzero elements in each row of the sparse matrix during the preprocessing stage on the CPU. For each row, the number of nonzeros is calculated by subtracting consecutive entries in the CSR row_ptr array. Based on a predefined threshold (e.g., 16 nonzeros), rows are categorized into two groups: low-density and high-density.

The indices of low-density rows are stored in a separate array (bin0_rows), while the indices of high-density rows are stored in another array (bin1_rows). These arrays are then transferred to GPU memory.

SpMV operation is divided into two specialized kernels: one is “spmv_bin0_kernel” and the other is “spmv_bin1_kernel”. During SpMV execution on the GPU, spmv_bin0_kernel uses bin0_rows to launch one thread per row for low-density rows, and spmv_bin1_kernel uses bin1_rows to launch one warp per row for high-density rows.

For spmv_bin0_kernel, rows with relatively few nonzeros (e.g., 16 or fewer) are assigned to spmv_bin0_kernel, where each row is processed by a single thread—this is efficient because the workload is small and thread-level parallelism suffices. The pseudo code for spmv_bin0_kernel is given below:

```

kernel spmv_bin0_kernel(row_ids, num_rows, row_ptr, col_ind, val, x, y):
    thread_id = blockIdx.x * blockDim.x + threadIdx.x

    if thread_id >= num_rows:
        return

    row = row_ids[thread_id]
    sum = 0

    for j from row_ptr[row] to row_ptr[row + 1] - 1:
        sum += val[j] * x[col_ind[j]]

    y[row] = sum

```

In contrast, rows with more nonzeros are assigned to spmv_bin1_kernel, where each row is handled by an entire warp (32 threads) benefiting the warp-level parallelism to compute partial sums in parallel. Then, partial sums are reduced efficiently using warp shuffle operations that do not require the use of shared memory.

The pseudo code for spmv_bin1_kernel is given below:

```

kernel spmv_bin1_kernel(row_ids, num_rows, row_ptr, col_ind, val, x, y):
    global_warp_id = (blockIdx.x * blockDim.x + threadIdx.x) / WARP_SIZE
    lane = threadIdx.x % WARP_SIZE

```

```

if global_warp_id >= num_rows:
    return

row = row_ids[global_warp_id]
sum = 0

// Each thread in the warp processes a subset of the row
for j from row_ptr[row] + lane to row_ptr[row + 1] - 1 step WARP_SIZE:
    sum += val[j] * x[col_ind[j]]

// Perform warp-wide reduction using shuffle instructions
for offset = WARP_SIZE / 2 down to 1 step /= 2:
    sum += warp_shuffle_down(sum, offset)

if lane == 0:
    y[row] = sum

```

Lastly, these two kernels are launched in separate CUDA streams to allow concurrent execution when hardware resources are sufficient. This selective kernel strategy makes sure better load balancing and efficient utilization of GPU threads across a wide range of sparse matrix structures.

D. Vector updates and dot products

Other CG operations such as vector updates and dot products are also parallelized across threads using standard 1D grid-block configurations. This approach ensures that all stages of the CG iteration can exploit the high parallelism offered by the GPU architecture. Additionally, dot product kernel uses shared memory and atomicAdd for safe accumulation across blocks. The pseudo code for dot product kernel is given below:

```

kernel vector_dot(a, b, result, n):
    shared temp[blockDim.x]
    tid = threadIdx.x
    i = blockIdx.x * blockDim.x + tid

    if i < n:
        temp[tid] = a[i] * b[i]
    else:
        temp[tid] = 0

    __syncthreads()
    // Parallel reduction in shared memory
    for offset = blockDim.x / 2 down to 1 step /= 2:
        if tid < offset:
            temp[tid] += temp[tid + offset]
        __syncthreads()

    if tid == 0:
        atomicAdd(result, temp[0])

```

E. Memory Usage

Memory usage in this implementation is designed to use the GPU's memory efficiently. All vectors and matrix data (in CSR format) are allocated in global memory, while shared memory is used in the dot product kernel to enable fast intra-block reductions.

Intermediate scalars used in convergence checks and coefficient calculations are kept in device memory and are accessed via cudaMemcpy only when necessary. Careful synchronization and memory allocation strategies help avoid race conditions and minimize latency.

V. EXPERIMENTS OR OTHER EVIDENCE OF SUCCESS

A. Parameter Settings

The following parameter settings were used to optimize performance and ensure correct convergence:

- **Maximum Iterations (MAX_ITEES):** Set to 10,000 to ensure sufficient opportunity for convergence, particularly for ill-conditioned matrices.
- **Tolerance (TOL):** A convergence threshold of $1e-11$ was used, terminating iterations when the residual norm falls below this value.
- **Thread Block Configuration:** For vector operations and spmv_bin0_kernel, a fixed block size of 256 threads was used. This setting was chosen based on typical occupancy and warp scheduling considerations on modern GPUs.
- **Warp-Based SpMV Configuration:** For spmv_bin1_kernel, 32-thread warps were grouped into blocks of 4 warps (i.e., THREADS_BIN1 = 128) to ensure balanced workload distribution and efficient warp utilization.
- **Binning Threshold:** Rows with 16 or fewer nonzeros were assigned to Bin0, while rows with more than 16 nonzeros were assigned to Bin1.
- **CUDA Streams:** Two streams were used to enable concurrent execution of spmv_bin0_kernel and spmv_bin1_kernel, improving throughput by overlapping execution when hardware resources allow.

B. Dataset

Sparse matrices used in this experiment were sourced from the Suite Sparse Matrix Collection [5] and represented in the Compressed Sparse Row (CSR) format due to its memory efficiency and suitability for SpMV operations. Additionally, a custom symmetric positive definite (SPD) matrix was generated in MATLAB to test balanced binning scenarios.

The dataset files, containing both the problem matrix and the right-hand side vector, were stored in .txt format for simplicity and ease of integration. The implementation reads four separate files: one for the Row Pointers, one for the Column Indices, one for the Values of the matrix, and one for the values of the right-hand side vector b.

C. Experimental Procedure

Two algorithmic approaches were used in the experiments: a CPU-based sequential Conjugate Gradient (CG) algorithm and a GPU-accelerated parallel CG implementation using CUDA. Both methods solve the same set of sparse linear systems using identical matrix and vector inputs to ensure fair comparison.

For each matrix, the total execution time of the CG solver was recorded. In the GPU implementation, CUDA event timers were used to measure kernel-level timings, including the combined time of spmv_bin0_kernel and spmv_bin1_kernel, as well as their individual execution times. The CPU execution time was recorded using wall-clock timing in C++.

To evaluate performance, speedup was computed as the ratio of CPU time to GPU time. Additionally, other metrics such as the number of CG iterations to converge and the row binning distribution (size of bin0 and bin1) were collected for each matrix. This procedure was repeated for all datasets, including standard benchmarks and the custom SPD matrix. The resulting data was then analyzed to understand the effect of matrix structure, kernel concurrency, and binning balance on performance.

D. Results

Tables II and III present the results of performance experiments comparing CPU and GPU implementations of the Conjugate Gradient (CG) algorithm on a range of sparse matrices. Table I provides structural information about each matrix, including the number of rows, columns, and nonzero entries, as well as the condition number where available. The condition number provides an indication of the expected convergence rate in the CG method—larger condition numbers generally lead to slower convergence and require more iterations. The right-hand side (RHS) vectors for all test cases were randomly generated in MATLAB to simulate general-purpose input. Table I also includes size(Bin0) and size(Bin1), which represent the number of rows assigned to two specialized SpMV kernels based on row sparsity—Bin0 for rows with fewer nonzeros (handled using a thread-per-row strategy), and Bin1 for denser rows (processed with a warp-per-row strategy). For entries marked N/A, the condition number was not specified or not available.

Table III summarizes the performance results, including the number of CG iterations until convergence, the total CPU execution time, and three GPU timing metrics. GPU Time 1 reports the total execution time when both SpMV kernels are launched concurrently in separate CUDA streams. GPU Time 2 and GPU Time 3 show the individual execution times of `spmv_bin0_kernel` and `spmv_bin1_kernel`, respectively, when measured in isolation. These results help assess the performance benefits of GPU acceleration, the effectiveness of the binning strategy, and the impact of concurrency on execution time across different matrix types.

TABLE I: Structural Information Table for Sparse Matrices

Matrix	Rows	Columns	Nonzeros	Cond. Num.	size(Bin0)	size(Bin1)
HB/bcsstm09	1083	1083	1,083	1.00E+04	1083	0
HB/nos3	960	960	15,844	3.77E+04	245	715
HB/bcsstk27	1224	1224	56,126	2.41E+04	4	1224
Pothen/bodyy5	18589	18589	128,853	7.87E+03	18589	0
MathWorks/Kuu	7102	7102	340,200	1.58E+04	0	7102
UTEP/Dubcova2	65025	65025	1,030,225	N/A	48900	16125
Custom SPD	30000	30000	4,543,292	N/A	14950	15050

In Table III, the results demonstrate the effectiveness and scalability of our GPU-accelerated Conjugate Gradient (CG) method across a diverse set of sparse matrices.

TABLE II: Speedup Table for Sparse Matrices

Matrix	CPU Time	GPU Time 1	Speedup (CPU/GPU1)
HB/bcsstm09	0.2864 ms	1.576 ms	0.18 ×
HB/nos3	47.0606 ms	88.385 ms	0.53 ×
HB/bcsstk27	514.721 ms	349.772 ms	1.47 ×
Pothen/bodyy5	1392.49 ms	208.902 ms	6.66 ×
MathWorks/Kuu	1702.81 ms	212.625 ms	8.01 ×
UTEP/Dubcova2	3484.62 ms	94.571 ms	36.85 ×
Custom SPD	612.47 ms	15.6725 ms	39.07 ×

TABLE III: Comparison of CPU and GPU Execution Times for Sparse Matrices

Matrix	Iteration	CPU Time	GPU Time 1	GPU Time 2	GPU Time 3
HB/bcsstm09	4	0.2864 ms	1.576 ms	1.437 ms	2.427 ms
HB/nos3	512	47.0606 ms	88.385 ms	101.563 ms	93.077 ms
HB/bcsstk27	2201	514.721 ms	349.772 ms	432.524 ms	352.226 ms
Pothen/bodyy5	1315	1392.49 ms	208.902 ms	210.487 ms	287.084 ms
MathWorks/Kuu	1151	1702.81 ms	212.625 ms	270.940 ms	210.663 ms
UTEP/Dubcova2	506	3484.62 ms	94.571 ms	90.035 ms	154.422 ms
Custom SPD	38	612.47 ms	15.6725 ms	268.503 ms	18.1784 ms

For small matrices such as HB/bcsstm09 and HB/nos3, the GPU implementation yields limited benefit due to the low computational intensity which is shown by the smaller execution time on the CPU than the one on the GPU. However, for larger and more complex matrices such as HB/bcsstk27, MathWorks/Kuu, and UTEP/Dubcova2 the GPU significantly outperforms the CPU. For example, in Dubcova2, it completes in 94.571 ms on the GPU versus 3484.62 ms on the CPU resulting in a speedup of 36.85×, even with over 2200 CG iterations. This gain is largely credited to efficient handling of dense rows via the warp-based `spmv_bin1_kernel`.

For bcsstk27, bodyy5, and Kuu, the GPU Time 1 closely matches the execution time of the dominant SpMV kernel. This is because the row binning is highly unbalanced in these matrices, nearly all rows fall into either bin0 or bin1. As a result, only one type of kernel is predominantly launched and the contribution of the other kernel is minimal. This skewed bin distribution leads to minimal overlap between the two kernels which is making the GPU Time 1 effectively determined by the heavier bin’s execution.

In the case of Dubcova2, although it is the largest matrix in the dataset, with over 65,000 rows and more than one million nonzeros, it achieves excellent GPU performance due to the binning structure and the low iteration count. Approximately 75% of the rows (48,900) fall into bin0, each having 16 or fewer nonzeros, and are processed by `spmv_bin0_kernel` using a thread-per-row strategy. The remaining 25% (16,125 rows) fall into bin1 where every one of these rows are processed by `spmv_bin0_kernel` using a warp-per-row strategy. Another reason for the exceptional performance in Dubcova2 is that it converges in just 506 iterations which are much fewer than those in matrices like bcsstk27 (which takes 2201 iterations) or nos3 (512 iterations). Since matrix-vector multiplications dominate the cost per iteration in CG, a low iteration count directly translates into faster overall convergence and lower run-

time. The combined execution of `spmv_bin0_kernel` and `spmv_bin1_kernel`, launched in separate CUDA streams, results in a total GPU time of 94.57 ms despite the individual kernel times being 90.03 ms and 154.42 ms respectively. This shows a stream-level concurrency where the kernels partially overlap in execution which is reducing the runtime. This stream-level concurrency between `spmv_bin0_kernel` and `spmv_bin1_kernel` is illustrated in Fig. 2 and Fig. 3 which are captured from the trace output of Nsight Systems 2025.1.3.




Fig. 2: Cropped timeline view of the GPU implementation.

#	Name	Start	Duration	GPU	Context
377	<code>spmv_bin0_kernel(const int *, int, const int *, const float *, const float *, float *)</code>	2.46077s	20.152 μ s	GPU 0	Stream 14
378	<code>spmv_bin1_kernel(const int *, int, const int *, const float *, const float *, float *)</code>	2.46078s	48.321 μ s	GPU 0	Stream 13
379	Memset	2.46080s	288 ns	GPU 0	Stream 7
380	<code>vector_dot(const float *, const float *, float *, int)</code>	2.46082s	7.168 μ s	GPU 0	Stream 7
381	Memory DtoH (Pageable)	2.46089s	736 ns	GPU 0	Stream 7
382	Memset	2.46091s	288 ns	GPU 0	Stream 13
383	<code>vector_sspv(int, float *, const float *, float)</code>	2.46091s	1.920 μ s	GPU 0	Stream 13
384	<code>vector_dot(const float *, const float *, float)</code>	2.46093s	1.888 μ s	GPU 0	Stream 14
385	<code>vector_dot(const float *, const float *, float *, int)</code>	2.46093s	7.072 μ s	GPU 0	Stream 13
386	Memory DtoH (Pageable)	2.46095s	736 ns	GPU 0	Stream 7
387	<code>vector_scale_add(int, float *, const float *, float)</code>	2.46097s	1.932 μ s	GPU 0	Stream 7
388	<code>spmv_bin0_kernel(const int *, int, const int *, const float *, const float *, float *)</code>	2.46098s	20.801 μ s	GPU 0	Stream 14
389	<code>spmv_bin1_kernel(const int *, int, const int *, const float *, const float *, float *)</code>	2.46099s	48.641 μ s	GPU 0	Stream 13
390	Memset	2.46707s	288 ns	GPU 0	Stream 7
391	<code>vector_dot(const float *, const float *, float *, int)</code>	2.46708s	7.200 μ s	GPU 0	Stream 7
392	Memory DtoH (Pageable)	2.46711s	768 ns	GPU 0	Stream 7
393	Memset	2.46712s	288 ns	GPU 0	Stream 13
394	<code>vector_sspv(int, float *, const float *, float)</code>	2.46712s	1.888 μ s	GPU 0	Stream 13
395	<code>vector_sspv(int, float *, const float *, float)</code>	2.46713s	1.888 μ s	GPU 0	Stream 14
396	<code>vector_dot(const float *, const float *, float *, int)</code>	2.46716s	7.104 μ s	GPU 0	Stream 13

Fig. 3: Cropped event view of the GPU implementation.

In Table III, it is observed that `spmv_bin0_kernel` alone appears faster than the combined execution of `spmv_bin0_kernel` + `spmv_bin1_kernel` in Dubcova2. The reason is that when measuring `spmv_bin0_kernel` in isolation, it runs without any other kernel occupying GPU resources. This allows it to use more of the GPU's compute and memory bandwidth. But when both kernels are launched together in separate streams, they compete for shared resources which may slightly delay the execution of bin0 internally even though the total time (combined) becomes shorter due to overlap. Another reason is that in bin1, each row contains exactly 25 nonzeros, which results in underutilization of GPU warps. Since each warp consists of 32 threads but only 25 elements need to be processed per row, the remaining 7 threads in each warp remain idle which is reducing the overall efficiency. In contrast, bin0 uses a thread-per-row strategy, which is better suited for the irregular but small row sizes in that bin and achieves higher utilization. In the case of nos3 where the rows of its bin1 have at least 32 nonzeros, the combined execution of `spmv_bin0_kernel` and `spmv_bin1_kernel` is actually faster than running each kernel alone. This indicates that the GPU has sufficient resources such as shared memory, execution units, and memory bandwidth to accommodate

both kernels concurrently without introducing significant delays. As a result, both kernels can execute efficiently in parallel while benefiting from overlap without degrading each other's performance.

To evaluate performance under balanced binning conditions, a custom symmetric positive definite (SPD) matrix was generated in MATLAB, where rows in bin0 each contain 10 nonzeros and rows in bin1 contain 64 nonzeros. The matrix was designed with nearly equal bin sizes and appropriate row density (14,950 rows with 10 nonzeros in bin0 and 15,050 rows with 64 nonzeros in bin1) to ensure a well-balanced workload between the two SpMV kernels. Due to this balance, both `spmv_bin0_kernel` and `spmv_bin1_kernel` are actively utilized and executed concurrently in separate CUDA streams. As a result, for Custom SPD in Table III, the total GPU execution time (GPU Time 1) is shorter than either kernel's individual runtime, clearly demonstrating the benefits of balanced row distribution and effective kernel overlap.

In conclusion, the results show that the GPU-based Conjugate Gradient (CG) method performs best when the matrix structure and workload distribution are well balanced. While small matrices don't gain much speedup due to low computation per iteration, larger matrices with many nonzeros and structured sparsity show significant improvements. Dividing matrix rows into two bins—one for rows with fewer nonzeros (handled by `spmv_bin0_kernel`) and one for denser rows (handled by `spmv_bin1_kernel`)—helps match the kernel strategy to the data pattern. When these two kernels are launched in separate CUDA streams, they can run at the same time, reducing overall execution time. However, this only works well if the GPU has enough resources and the row sizes are suited to the kernel. For example, in Dubcova2, rows in bin1 have exactly 25 nonzeros, which leads to underuse of GPU threads in a warp, slightly lowering performance. Still, thanks to a high percentage of simple rows in bin0 and a low number of CG iterations, Dubcova2 runs very efficiently on the GPU. A custom balanced matrix also confirms that when the workload between the two bins is even, both kernels can overlap effectively, resulting in better overall performance. This highlights the importance of load balancing and stream concurrency when optimizing CG on GPUs.

VI. DISCUSSION AND FUTURE WORK

The results of this study highlight the effectiveness of stream-level concurrency and bin-based kernel decomposition in accelerating the Conjugate Gradient (CG) method on GPUs. By dividing matrix rows into two categories based on sparsity and assigning them to specialized kernels (`spmv_bin0_kernel` and `spmv_bin1_kernel`), the implementation achieves better load balancing and

improved resource utilization. Launching these kernels in separate CUDA streams enables concurrent execution, which, when combined with a balanced bin distribution, leads to significant reductions in total execution time. This effect is particularly pronounced in matrices with a large number of rows and structured sparsity, such as Dubcova2 and the custom SPD matrix, where speedups exceeded 30x over CPU execution.

However, several limitations are evident. The use of a fixed binning threshold does not generalize well across all matrices, and resource contention between concurrent kernels can sometimes reduce the efficiency of individual kernels, even if the total time improves. Additionally, the implementation lacks support for preconditioning, which is often essential for solving ill-conditioned systems efficiently. The design is also limited to single-GPU environments and does not address scalability across multiple GPUs or distributed systems. Despite these limitations, the proposed method is well-suited for large-scale, symmetric positive definite problems common in scientific simulations and engineering applications.

Future work should focus on adaptive or multi-bin strategies for more flexible row classification, integration of GPU-friendly preconditioners, and extension to multi-GPU or hybrid CPU-GPU architectures. Further enhancements such as auto-tuning of kernel configurations and benchmarking against established libraries like cuSPARSE would also provide deeper insight into the competitiveness and generalizability of the proposed approach.

REFERENCES

- [1] A. E. Gogusdere, "CENG577 - Final Project Report," METU, Ankara, 2025.
- [2] P. Ghysels and W. Vanroose, "Hiding global synchronization latency in the preconditioned Conjugate Gradient algorithm," *Parallel Computing*, vol. 40, no. 7, pp. 224–238, Jun. 2013, doi: 10.1016/j.parco.2013.06.001.
- [3] J. Cornelis, S. Cools, and W. Vanroose, "The Communication-Hiding Conjugate Gradient Method with Deep Pipelines," *arXiv.org*, Jan. 15, 2018. <https://arxiv.org/abs/1801.04728>
- [4] A. T. Chronopoulos and C. W. Gear, "s-step iterative methods for symmetric linear systems," *Journal of Computational and Applied Mathematics*, vol. 25, no. 2, pp. 153–168, Feb. 1989, doi: 10.1016/0377-0427(89)90045-9.
- [5] "SuiteSparse Matrix Collection," the University of Florida, [Online]. Available: <https://sparse.tamu.edu/>. [Accessed 2025].